



# AutoVision

XLR8 PROJECT - ITSG 2025-2026

PRESENTERS: DAVID SZILAGYI, RAZVAN FILEA,  
ARMIN TOROK

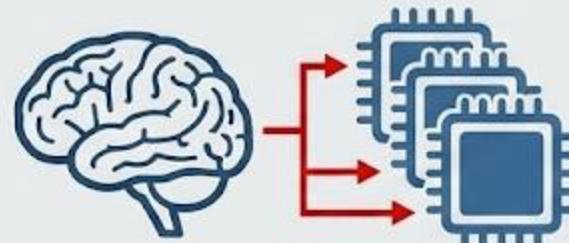
# THE CHALLENGE: REAL-TIME INTELLIGENT PERCEPTION

## THE CONTEXT



Simultaneous detection,  
tracking, recognition.

## THE PROBLEM



DEEP LEARNING  
(CNNs) COMPUTATIONALLY EXPENSIVE



SYNCHRONIZATION  
DELAYS & OVERHEAD

## THE TRADE-OFF



RAPID  
PROTOTYPING  
(Python)

DEPLOYMENT  
EFFICIENCY  
(C++/Rust)

# Initial ML Approach: Fine-Tuning on Real-World Data (BDD100K)

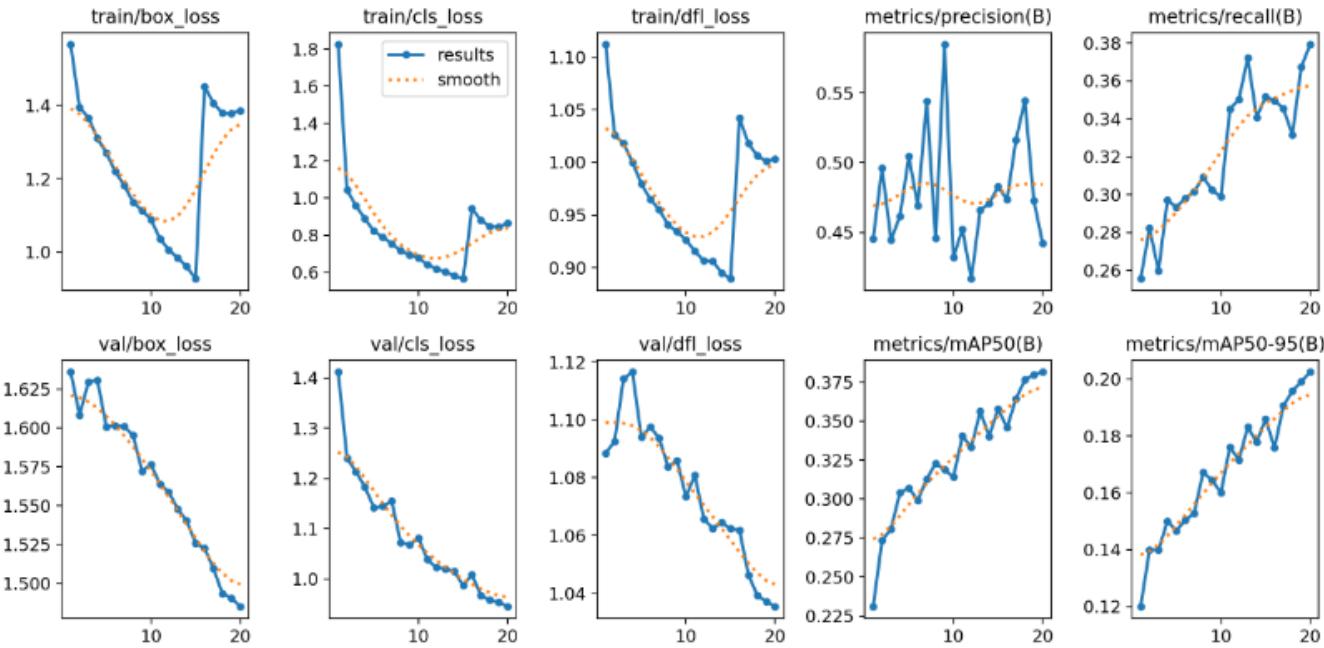


Figure 5.2: Training and validation loss evolution along with precision, recall, and mAP metrics across epochs. The late-epoch fluctuation corresponds to mosaic augmentation being disabled.

- Methodology:
  - Fine-tuned YOLO 11s (COCO-pretrained) on the Berkeley DeepDrive 100K dataset.
  - Training: 20 epochs, AdamW optimizer,  $640 \times 640$  resolution, focused on traffic lights and signs.
- Theoretical Performance:
  - Achieved stable convergence with validation  $mAP@0.5$  of 0.38 .
  - High precision for vehicles (0.90) and moderate precision for lights/signs in urban images. (0.53, 0.56)
- The "Sim-to-Real" Failure:
  - Critical Domain Gap: The model failed when deployed in the 1:10 scale in-office environment.
  - False Positives: Background objects (chairs, fire extinguishers) were frequently misclassified as traffic signals .

# Expectation meets Reality: The Sim-to-Real Gap

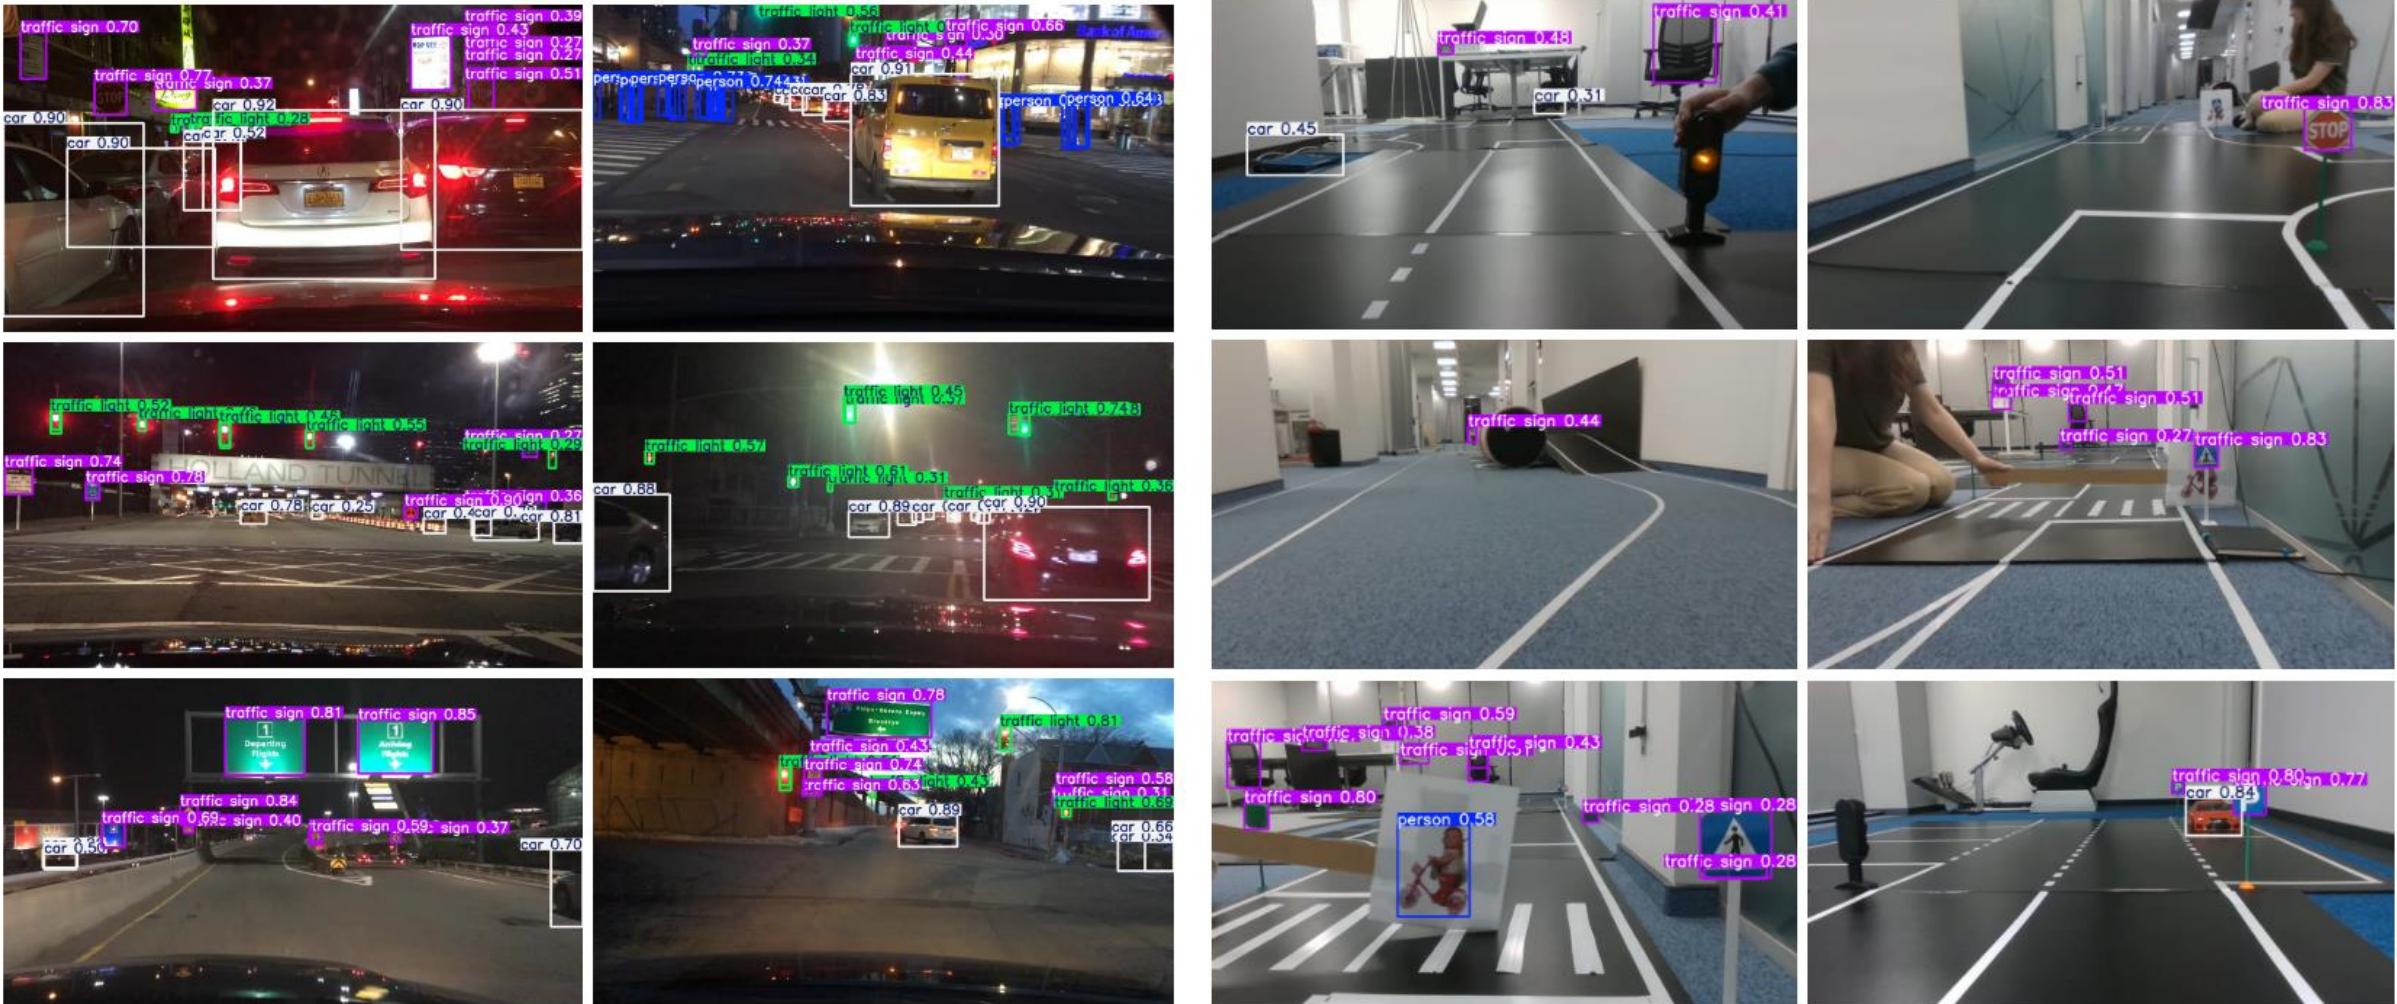
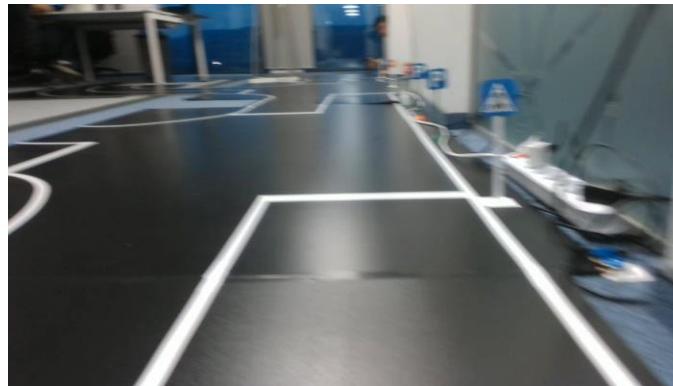


Figure 5.6: Qualitative detection results on the BDD100K test set using the fine-tuned YOLO11s model.

Figure 5.7: Detection performance in the real test environment.

# The Pivot to Sensor-Native Data

- Data Strategy:
  - Abandoned large-scale external datasets in favor of a **Handcrafted Dataset** captured directly from the vehicle's onboard sensors.
  - **Data Collection:** Manual (RC-mode) driving sessions to capture exact lighting and camera angles.
- Data Pipeline:
  - Applied custom augmentations: Brightness jitter, Gaussian noise, and rotation to simulate sensor noise and diversity .



# Model Specialization

- Architectural Change:
  - Replaced the monolithic model with **three specialized YOLO models + a geometric Lane Detector**
  - Why Specialization?
    - **Direct Classification:** Detects states (e.g., Red/Green) immediately, eliminating the need for a secondary classification stage
    - **No Compromises:** Avoids trade-offs in resolution and anchors by independently tuning for disparate objects
- Key Result:
  - Significantly mitigated background false positives and achieved robust detection of miniature props in real-time.

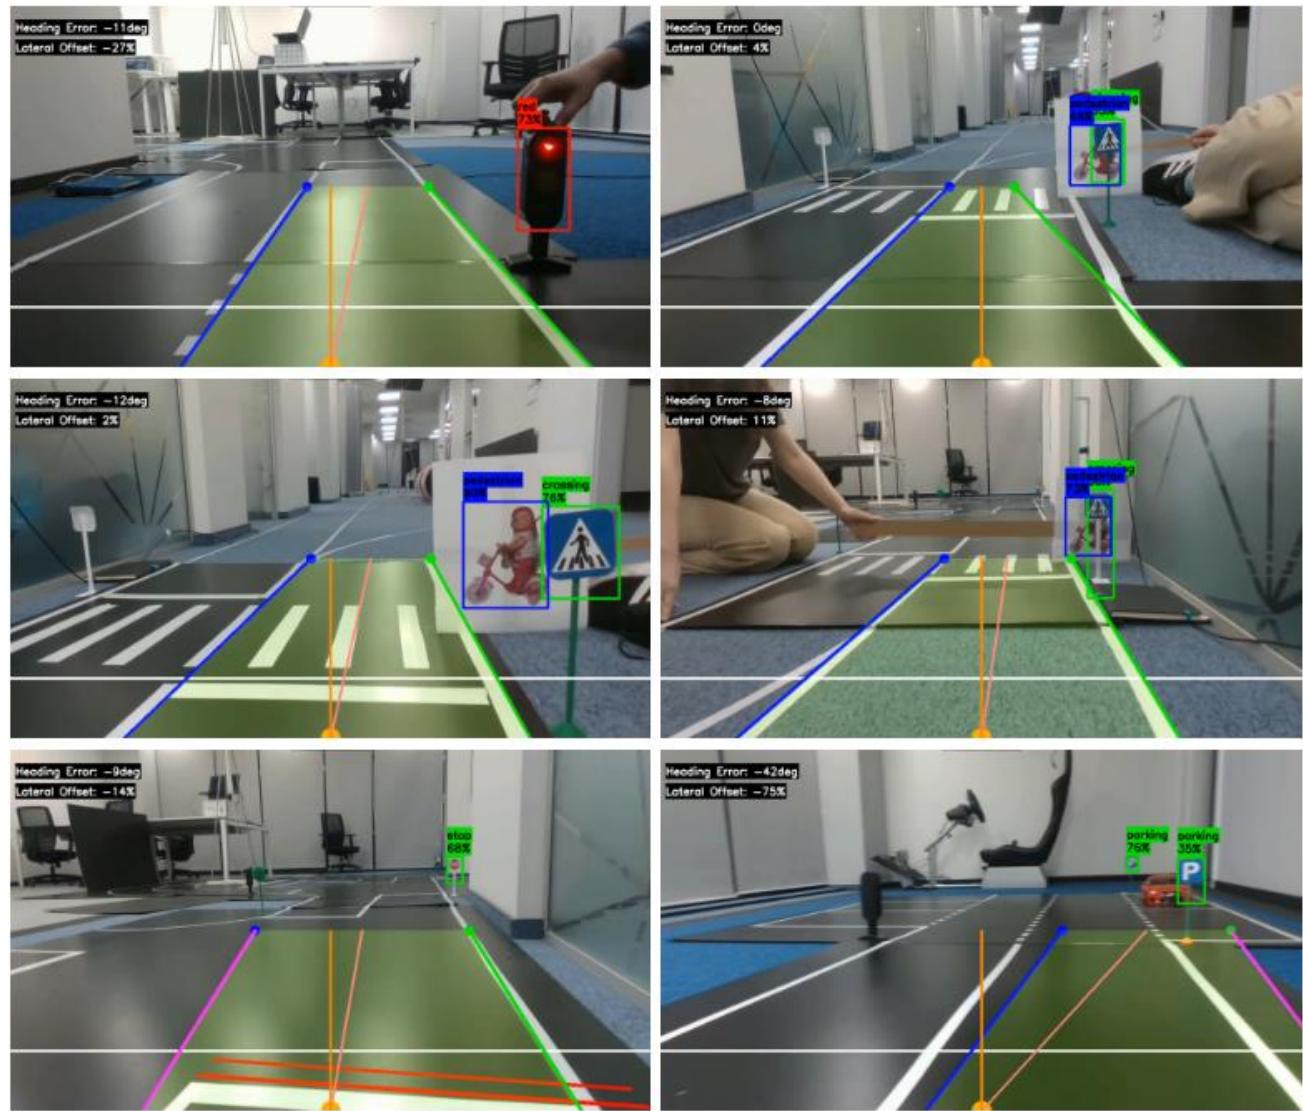
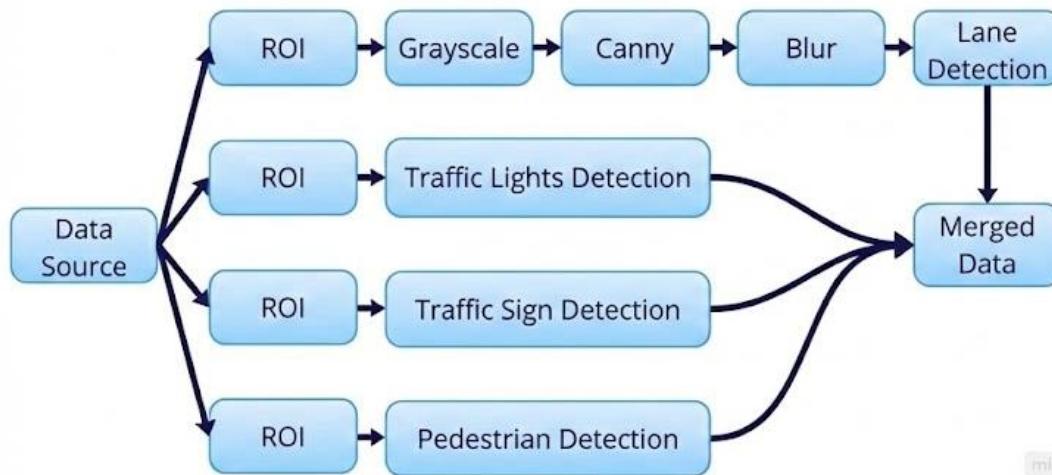
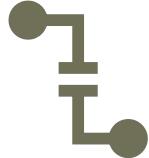


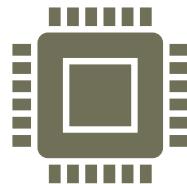
Figure 5.9: Qualitative detections in the scaled test environment after pretraining on the handcrafted dataset.

# The New Bottleneck: Concurrency & Latency



## The Consequence of Specialization:

We moved from 1 model to **4 concurrent perception branches**, forcing parallelism for performance



## Python's Limitation:

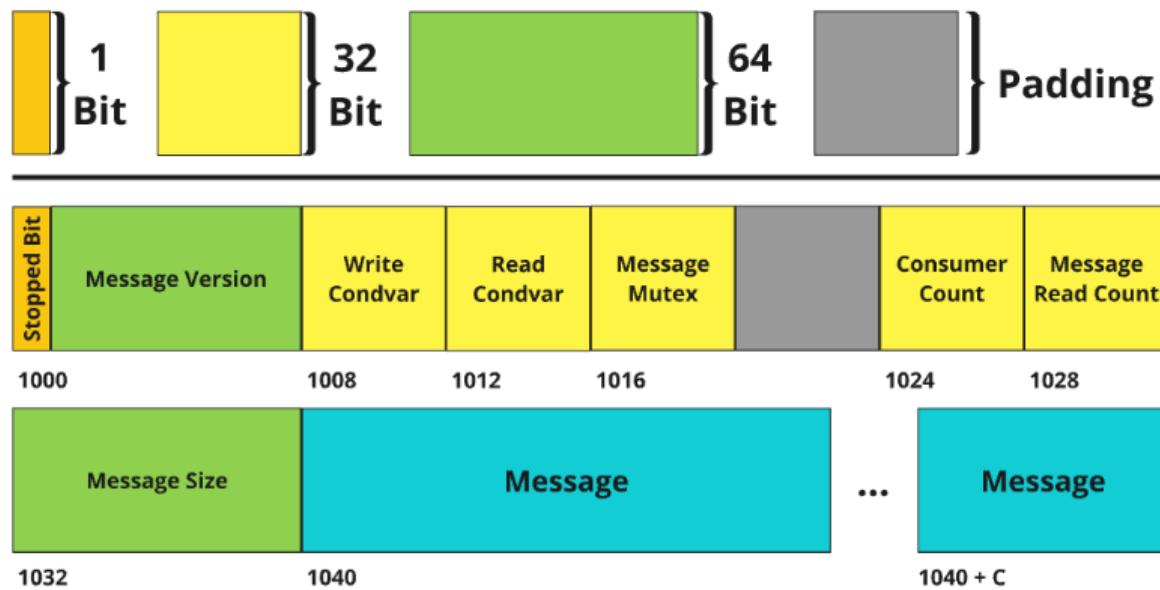
Standard Python multiprocessing struggles with **high-throughput video streams**.  
Leads to **high latency and memory overhead**.



## The Timing Requirement:

Fast branches (ex Lane Detection, ~6ms) must **update continuously** for steering control.  
Slow branches (ex Sign Recognition, ~28ms) must **not stall** the entire system.

# The Shared-Message Specification



**Table 1.** Write-Read Average Transfer Time (ms) 50.000 iterations

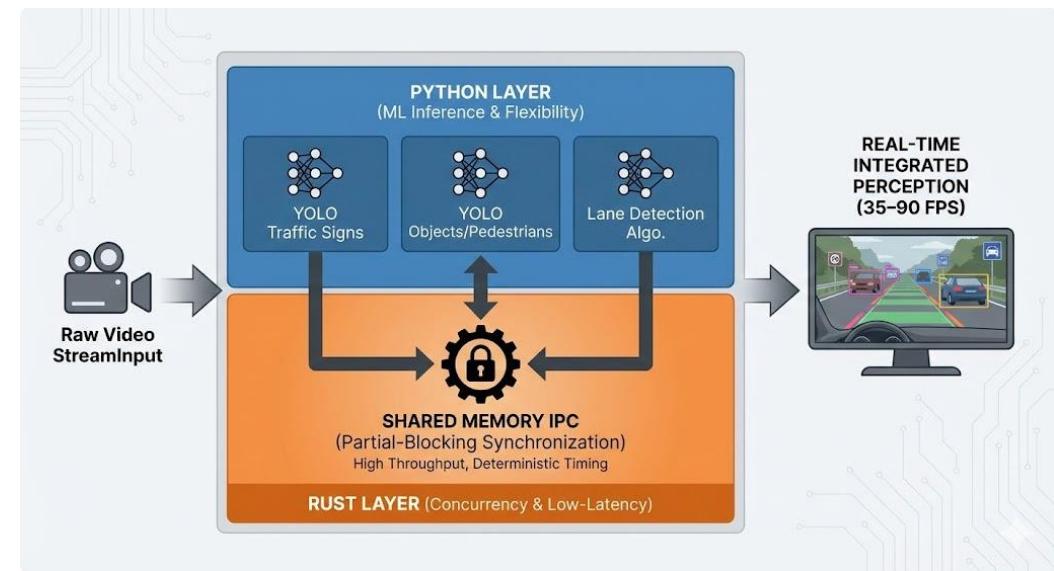
| Resolution | Size     | MP-Pipe | RS-IPC | SHM-LOCK | SHM-PTH |
|------------|----------|---------|--------|----------|---------|
| 256x256    | 576 KB   | 0.247   | 0.2    | 0.226    | 0.233   |
| 512x512    | 2.25 MB  | 0.908   | 0.728  | 0.796    | 0.82    |
| 640x480    | 2.64 MB  | 1.019   | 0.898  | 0.982    | 0.994   |
| 1280x720   | 7.91 MB  | 5.542   | 4.529  | 4.762    | 4.833   |
| 1920x1080  | 17.80 MB | 16.161  | 11.26  | 11.35    | 11.96   |

**Table 2.** Write-Read Total Transfer Time (s) 50.000 iterations

| Resolution | Size     | MP Pipe | RS-IPC  | SHM-LOCK | SHM-PTH |
|------------|----------|---------|---------|----------|---------|
| 256x256    | 576 KB   | 12.336  | 10.01   | 11.306   | 11.67   |
| 512x512    | 2.25 MB  | 45.411  | 36.408  | 39.789   | 40.986  |
| 640x480    | 2.64 MB  | 50.931  | 44.895  | 49.12    | 49.986  |
| 1280x720   | 7.91 MB  | 277.103 | 226.434 | 238.106  | 241.671 |
| 1920x1080  | 17.80 MB | 808.053 | 562.98  | 567.252  | 598.005 |

# Proposed Solution - Hybrid Architecture

- Concept: A "SharedMessage" infrastructure using Rust for low-level memory management
- Design:
  - Python Layer: Handles ML inference logging.
  - Rust Layer: Handles atomic synchronization, shared memory, and data flow and more
- Benefit: Enables high throughput without sacrificing the simplicity of Python development.



# The Core Algorithm - Partial-Blocking IPC

## The Communication Model.

- A Producer-Consumer strategy tailored for robotics.

## The Innovation: "Partial-Blocking" policy.

- Producer (Camera): Advances to the next frame once at least one consumer started processing it

## Why "Partial-Blocking"?

Non-Blocking: Risks Desynchronization (*input mismatch*)

Full-Blocking: Slowest model bottlenecks whole pipeline

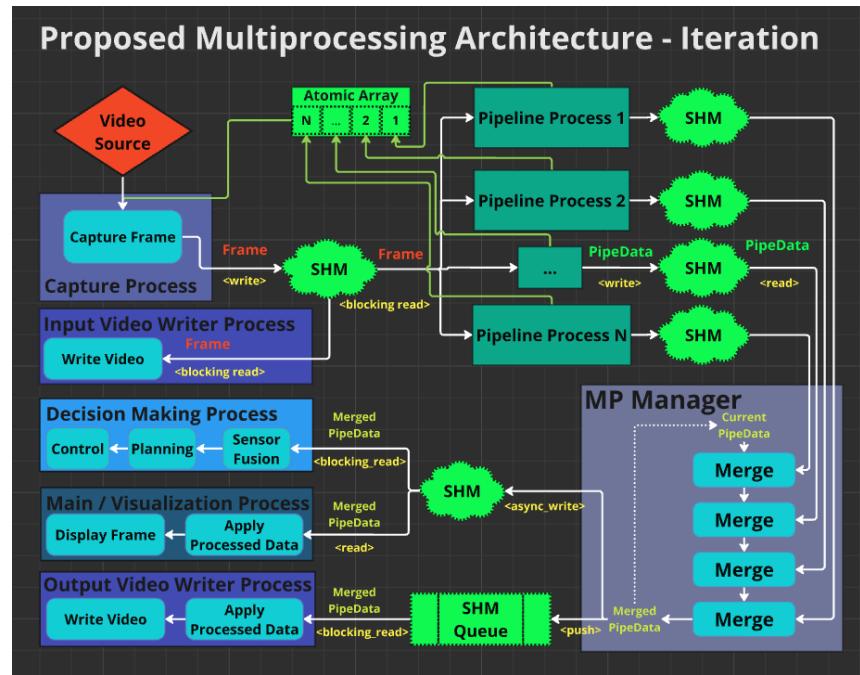
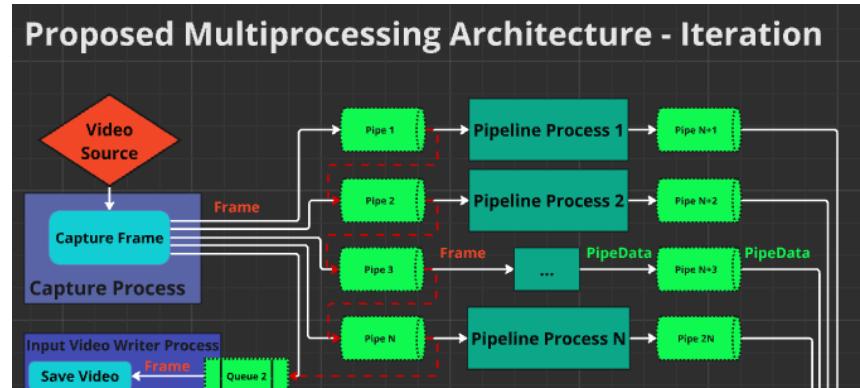
Partial-Blocking: Ensures frame consistency while allowing fast branches to run at maximum speed

## The Outcome.

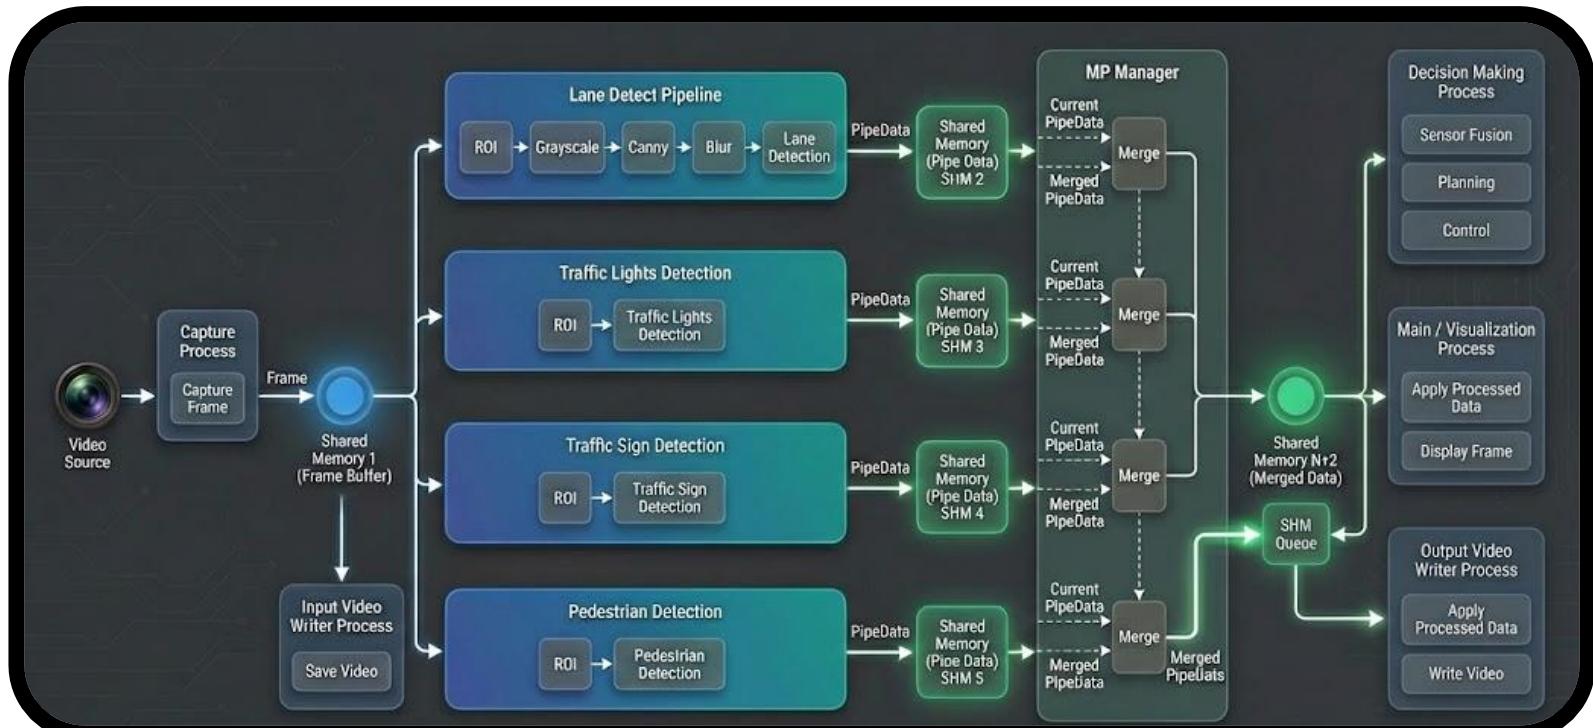
- Fast branches (lane detection) update frequently; slow branches (sign recognition) catch up asynchronously without stalling the system.

# Hybrid Python-Rust Architecture

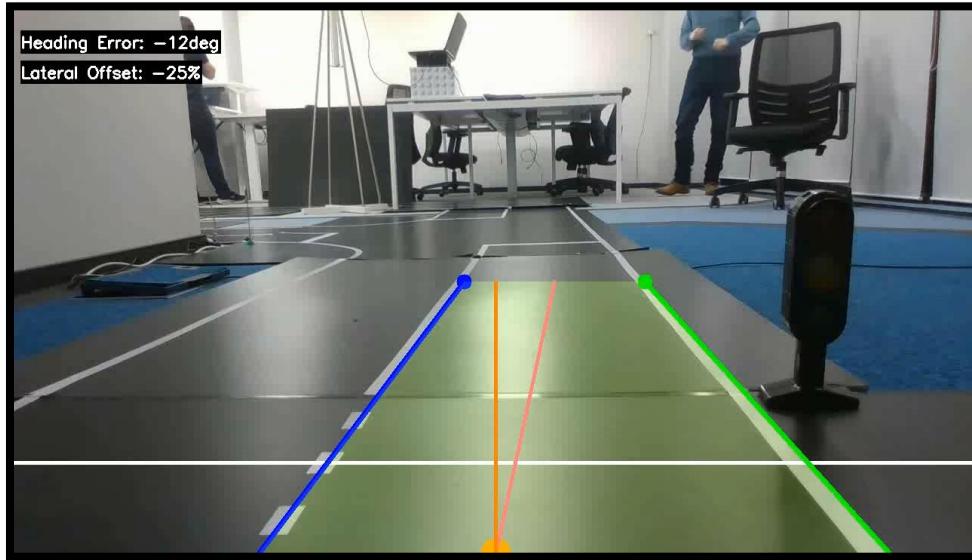
- **Design Philosophy:**
    - **Python:** Retained for ML development, model loading, and visualization.
    - **Rust:** Handles low-level memory management, atomic synchronization, and data flow.
  - **Mechanism:**
    - Modules communicate via a **Shared-Memory (IPC)** layer.
    - Ensures zero-copy data access where possible to maximize throughput.
  - **Benefit:**
    - Allows models to be retrained/swapped in Python without modifying the compiled Rust infrastructure .



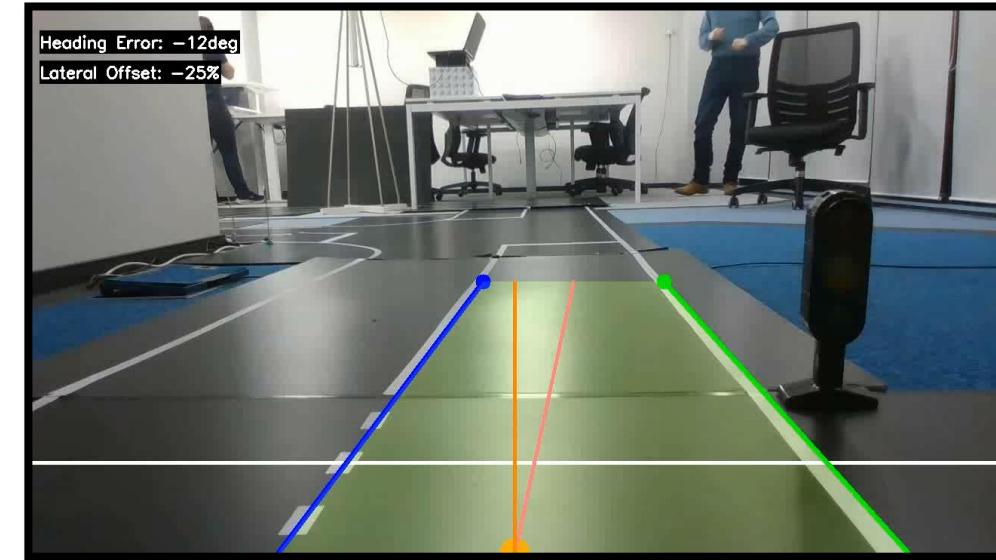
# DEMO



## **Full-blocking Approach With standard IPC (~20FPS)**

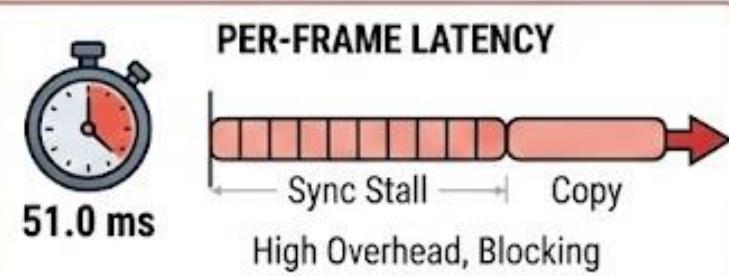
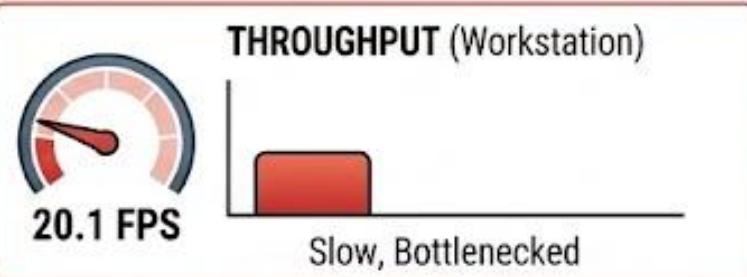


## **Proposed Shared-Memory Partial-Blocking Architecture (~90FPS)**

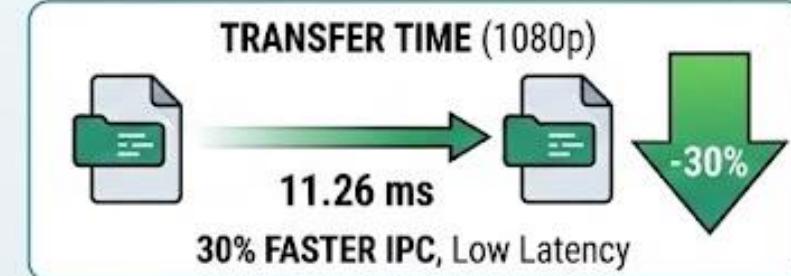
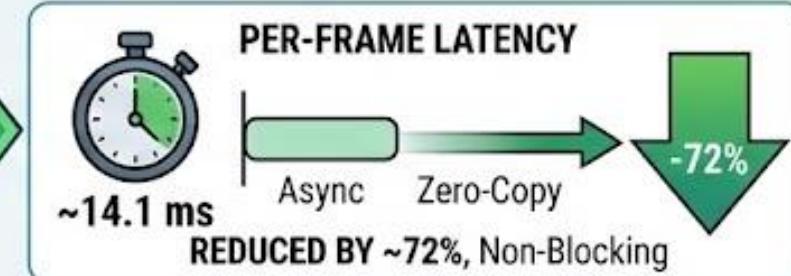


# QUANTITATIVE RESULTS: PERFORMANCE VALIDATION DIAGRAM

## NAIVE MULTIPROCESSING (Baseline)

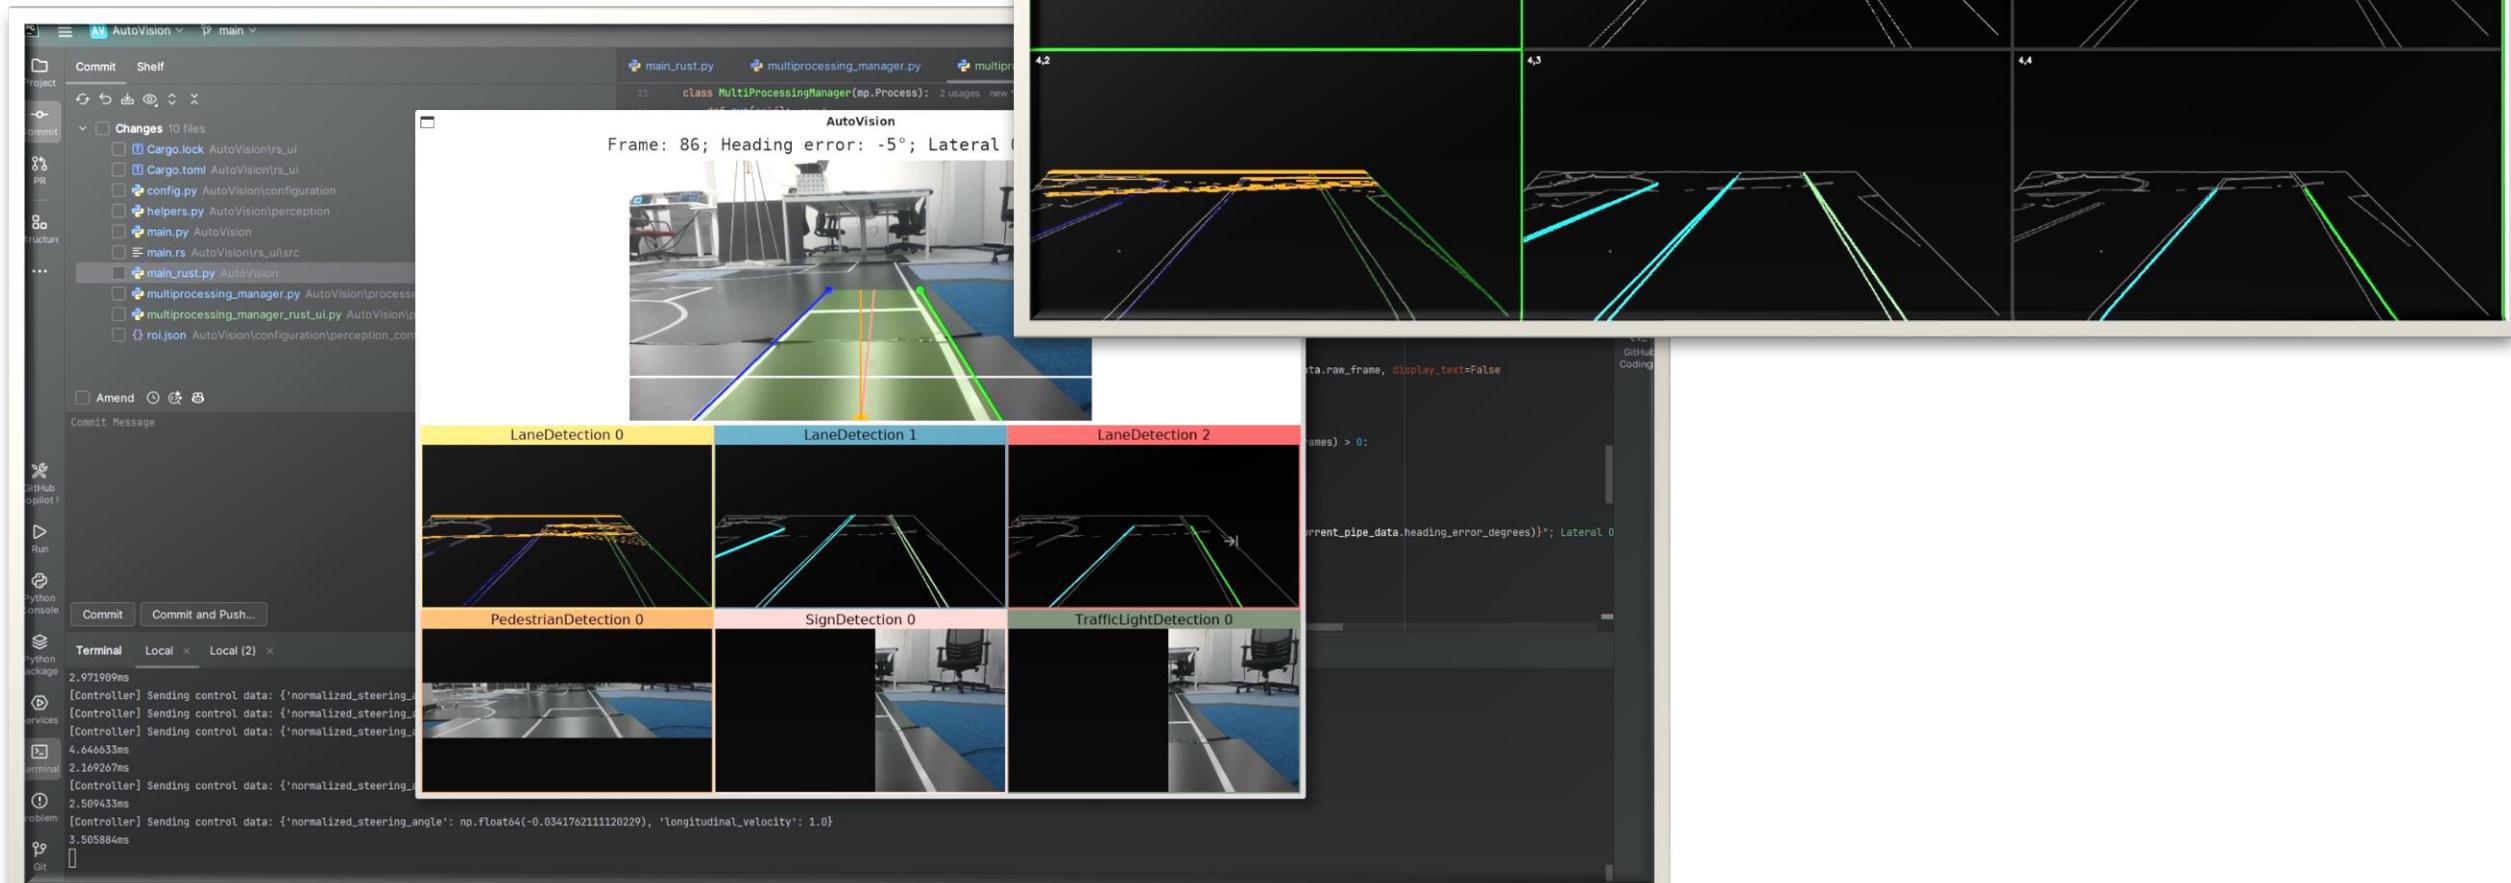


## HYBRID PYTHON-RUST (Ours)

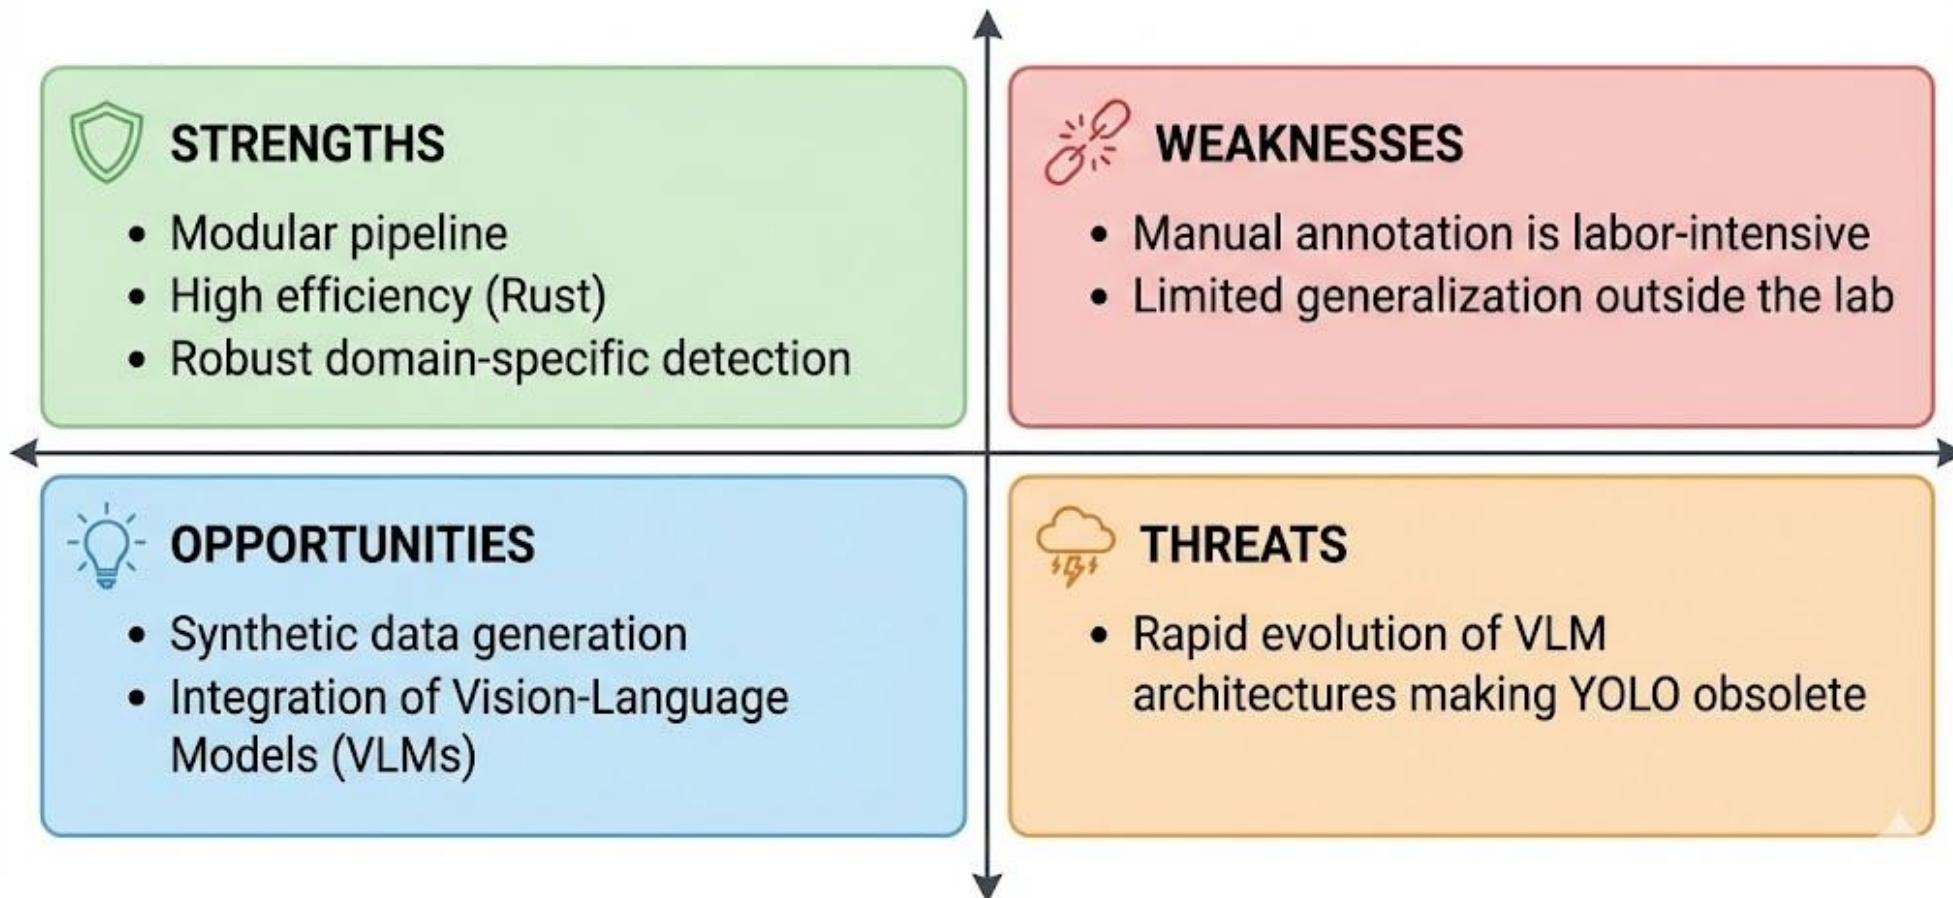


✓ CONCLUSION: Real-Time Performance (35+ FPS) Achieved on Embedded Hardware (Jetson AGX Orin)

# New GUI



# SWOT ANALYSIS: HYBRID PERCEPTION ARCHITECTURE



# Conclusion & Future Work

- Summary:
  - We introduced a modular architecture combining Python's flexibility with Rust's safety and efficiency.
  - The Partial-Blocking strategy successfully balances throughput and consistency.
- Impact:
  - Validated real-time performance on embedded devices, providing a scalable foundation for robotics research
  - Companion Paper published in [2025 RRIA, Romanian Journal of Information Technology and Automatic Control](#)
- Future Work:
  - Next-Gen Perception: Future iterations will explore integrating Vision-Language Models (VLMs) for semantic reasoning and adapting the runtime for the complex control stacks of Humanoid Robots.



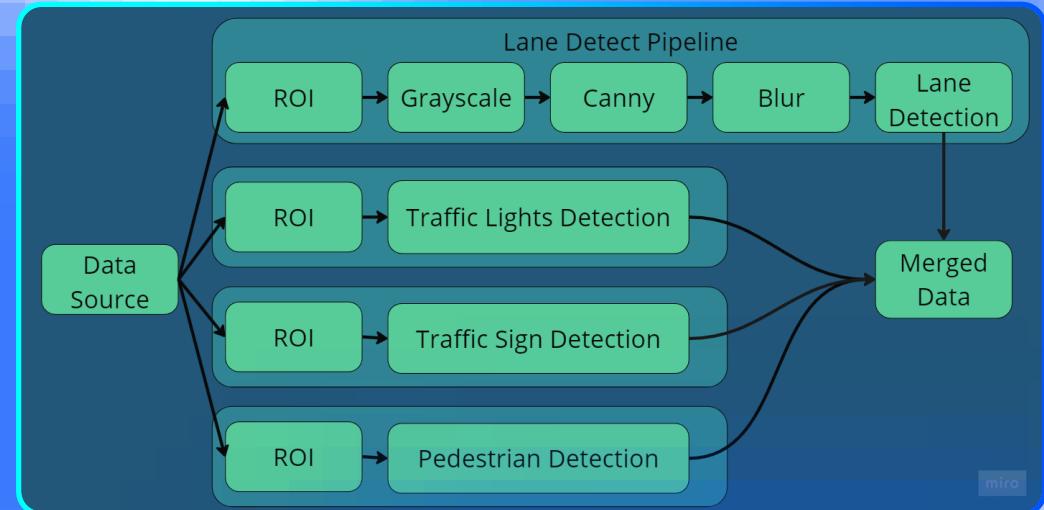


# THANK YOU

Questions? | XLR8 Project

# Perception System Multi-Processing Pipeline

```
{  
    "roi": {  
        "roi_type": "signs",  
        "visualize": false  
    },  
    "signs_detect": {  
        "model": "best_signs_close.pt",  
        "visualize": true  
    }  
},  
{  
    "roi": {  
        "roi_type": "traffic_lights",  
        "visualize": false  
    },  
    "traffic_light_detect": {  
        "model": "best_lights.pt",  
        "visualize": true  
    }  
},  
{  
    "roi": {  
        "roi_type": "pedestrians",  
        "visualize": false  
    },  
    "pedestrian_detect": {  
        "model": "best_pedestrian.pt",  
        "visualize": true  
    }  
}  
,  
{  
    "grayscale": {  
        "visualize": false  
    },  
    "canny_edge": {  
        "low_threshold": 150,  
        "high_threshold": 250,  
        "visualize": false  
    },  
    "blur": {  
        "kernel_size": 7,  
        "sigmaX": 1,  
        "visualize": false  
    },  
    "dilation": {  
        "kernel_size": 1,  
        "iterations": 1,  
        "visualize": false  
    },  
    "roi": {  
        "roi_type": "lines",  
        "visualize": true  
    },  
    "lane_detect": {  
        "visualize": false  
    },  
    "heading_error": {  
        "visualize": false  
    }  
}
```

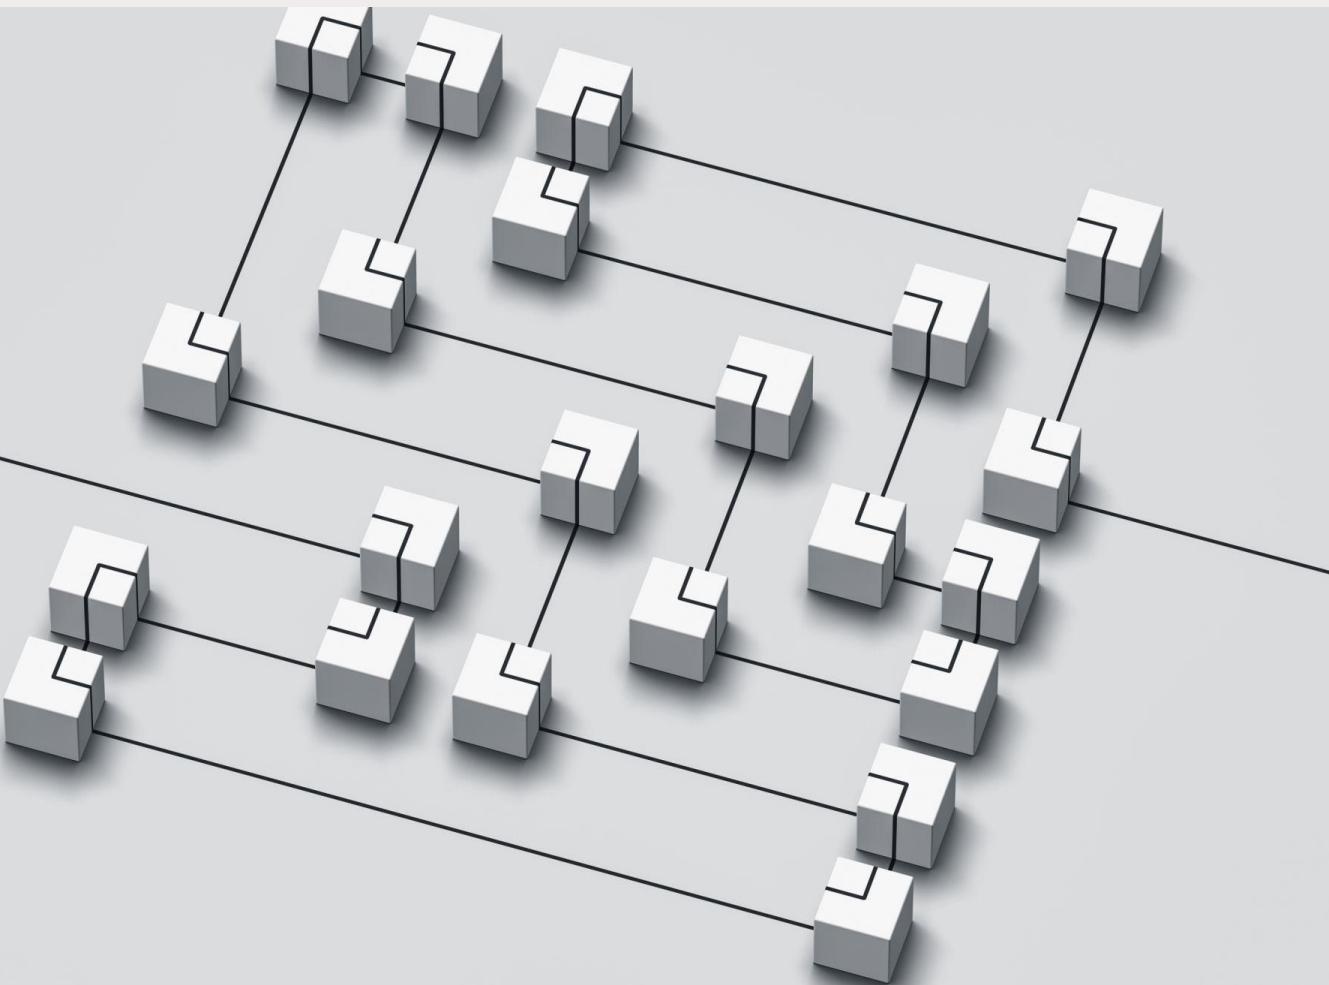


**Modular & Scalable Data Processing Pipeline**

**JSON Based Configuration of Filters and their Positioning**

**Ideal for Parallelized Workflows such as Autonomous Driving Perception Systems**

# Performance Validation



- **Throughput:**
  - The Shared-Memory architecture achieved **86.9 FPS** on a standard workstation, a **4.3x speedup** over the naive approach (20.1 FPS).
  - Achieved real-time performance (**35+ FPS**) even on embedded hardware (Jetson AGX Orin).
- **Latency Reduction:**
  - Total per-frame iteration time reduced by **~72%** by eliminating synchronization bottlenecks
  - Raw IPC transfer time for 1080p frames improved by **30%** ( $16.16\text{ms} \rightarrow 11.26\text{ms}$ ).

# The Challenge: Real-Time Intelligent Perception

- **The Context:**

- Autonomous driving requires simultaneous lane detection, object tracking, sign recognition and more.
- Safety depends on real-time responsiveness; a slow accurate detector is dangerous.

- **The Problem:**

- Deep Learning (CNNs) is computationally expensive.
- Executing multiple ML models concurrently introduces synchronization delays and overhead.

- **The Trade-off:**

- Researchers must choose between **Rapid Prototyping** (Python) and **Deployment Efficiency** (C++/Rust) .